

# **Lab Week 14 – CNN Image Classification**

Name: Pinisetti Sudhiksha

SRN: PES2UG23CS916

Course: Machine Learning

Submission Date: 19/11/2025

Section: F

## **Introduction**

The objective of this lab was to design, implement, and train a Convolutional Neural Network (CNN) capable of classifying images of hand gestures into three categories: rock, paper, and scissors. Using PyTorch and a labeled dataset, the task involved preparing the data, defining an appropriate CNN architecture, training the model, and evaluating its performance on unseen images.

## **Model Architecture**

The CNN architecture consisted of three convolutional blocks, each containing a  $3 \times 3$  convolutional layer followed by a ReLU activation and a MaxPool2d layer with a pooling size of 2. The network increased its channel depth from 3 to 16, then 32, and finally 64 across the convolutional stages while progressively reducing the spatial dimensions from  $128 \times 128$  down to  $16 \times 16$ . After the convolutional feature extractor, the model used a fully connected classifier that first flattened the  $64 \times 16 \times 16$  tensor into a vector of 16,384 elements. This was passed through a linear layer of size 256, followed by ReLU activation and a dropout layer with  $p=0.3$ , before reaching the final linear layer that outputs class logits for the three gesture categories.

## **Training and Performance**

Key hyperparameters:

- Optimizer: Adam
- Learning rate: 0.001
- Loss function: CrossEntropyLoss
- Epochs: 10

Performance results:

- Training loss reduced steadily from 0.7475 → 0.0029 across 10 epochs
- Final test accuracy: 97.95%

## **Conclusion and Analysis**

The model demonstrated strong performance, achieving nearly 98% accuracy on unseen data.

Loss decreased smoothly during training, indicating stable learning.

Individual image tests and example predictions confirmed that the model made correct classifications.

A few potential improvements include:

- Adding data augmentation to improve robustness
- Increasing model depth or experimenting with different architectures like BatchNorm to enhance generalization